

AI Deepfake Analyzer Report

The image has signals consistent with AI-generated or synthetic media.

Verdict: likely_ai_generated with medium confidence. Estimated AI probability is 79%; manipulation probability is 25%. Review the evidence layers before taking action.

Verdict	likely_ai_generated
Confidence	medium
AI Probability	79%
Manipulation Probability	25%
Detector Disagreement	64%

Decision Summary

The AI-generated verdict is mainly driven by visual detector evidence, then checked against metadata and forensic layers.

Primary drivers	Visual Model Consensus Layer: supports_ai_generated (AI 79%, manipulation 25%).
Counter-evidence	Compression And ELA Layer: supports_real_or_camera_origin (AI 2%, manipulation 24%).
Uncertainty	Detector disagreement is elevated at 64%. No signed C2PA/content credential was available to independently confirm provenance. 2 layer(s) produced weak anomaly signals that need context. 6 layer(s) were neutral or inconclusive.
What would help	Original, uncompressed media from the device or platform export. Public source URL, timestamps, captions, and account/page context. Signed C2PA/content credentials where available. Manual review when the result affects safety, reputation, or legal action.

Calibration Gate

Policy: victim_safe_calibrated_triage

Likely AI requires	Readable generative metadata or C2PA declaration,; or independent non-model fo
False-positive controls	Raw model scores are reliability-weighted before aggregation,; Detector disagree
Input quality risk	low

Mixture Of Experts

Expert	Opinion	Stance	Score
Visual Detector Ensemble	ensemble_likely_ai_generated	supports_ai	79%
Portrait Specialist	portrait_gate_not_met	neutral	
Metadata And Provenance	no_generative_marker	neutral	50%
Forensic Residuals	neutral	neutral	24%
Input Quality Guard	low_quality_risk	neutral	0%

Safety Arbiter	likely_ai_generated	supports_ai	79%
----------------	---------------------	-------------	-----

Evidence Layers

Input and Source - complete

- Image dimensions: 1254 x 1254.
- Input source: direct upload.

Metadata - complete

- No EXIF metadata found.
- No GPS metadata was exposed.
- No software marker found.

Provenance - unavailable

- No c2patool/c2pa executable was found on PATH.

Hashes - complete

- SHA-256: 7a3b93227153a932c461e92f7ec0a4ca573fced1cc960b5c8167793bd442ab96.
- Perceptual hash: 3fdf7e2602d99006.

Compression and Noise - complete

- Manipulation heuristic score: 0.2429.
- Artificiality heuristic score: 0.0247.

Detector Ensemble - complete

- metadata_provenance: inconclusive (ok).
- compression_noise_forensics: no_strong_forensic_signal (ok).
- hf:Ateeqq/ai-vs-human-image-detector: model_likely_ai_generated (ok).
- hf:dima806/ai_vs_real_image_detection: model_inconclusive (ok).
- hf:jacoballessio/ai-image-detect-distilled: model_likely_ai_generated (ok).

Explainable AI Trace

- Final label is likely_ai_generated with medium confidence.
- Combined AI probability is 79%.
- Combined manipulation probability is 25%.
- Layer ledger: 1 support AI, 1 support real/camera-origin, 0 support manipulation, 2 weak anomalies, 6 neutral/inconclusive.

Region Evidence Map

This abstract map compares regions against the rest of the same image. It does not display the uploaded image and does not identify what caused a regional difference.

R1 C1 31% low	R1 C2 31% low	R1 C3 38% low	R1 C4 47% medium
R2 C1 38% low	R2 C2 17% low	R2 C3 22% low	R2 C4 37% low
R3 C1 37% low	R3 C2 100% high	R3 C3 68% medium	R3 C4 33% low
R4 C1 32% low	R4 C2 76% high	R4 C3 69% medium	R4 C4 33% low

Analytical Layer Breakdown

Layer	Conclusion	AI	Manip.
Source, Metadata, And Provenance	neutral_or_inconclusive	50%	25%
Input Quality And Robustness	neutral_or_inconclusive	50%	0%
Visual Model Consensus Layer	supports_ai_generated	79%	25%
Luminance Layer	neutral_or_inconclusive	36%	25%
Chroma And Color Layer	neutral_or_inconclusive	38%	25%
Edge And Geometry Layer	neutral_or_inconclusive	39%	32%
Noise Residual Layer	neutral_or_inconclusive	35%	48%
Compression And ELA Layer	supports_real_or_camera_origin	2%	24%
Frequency Spectrum Layer	weak_anomaly	62%	43%
Tile Region Anomaly Layer	weak_anomaly	49%	56%

Source, Metadata, And Provenance: EXIF/XMP/PNG text extraction plus optional C2PA manifest inspection.

- Generative markers: none.
- Editing markers: none.

Input Quality And Robustness: Check resolution, file size, aspect ratio, entropy, and format conditions known to affect detector reliability.

- No major input-quality risk flags were detected.
- Resolution: 1254 x 1254; shortest edge 1254px.

Visual Model Consensus Layer: Run configured Hugging Face image-classification detectors and aggregate votes.

- 2/3 models voted AI-generated.
- 0/3 models voted real/human-origin.

Luminance Layer: Analyze grayscale entropy, contrast, clipping, and local brightness variation.

- Grayscale entropy: 7.435.
- Luminance contrast standard deviation: 0.190.

Chroma And Color Layer: Analyze RGB channel spread, channel correlation, and saturation distribution.

- Mean saturation: 0.577.
- Saturation variation: 0.232.

Edge And Geometry Layer: Run edge extraction and compare edge density/variance across tiles.

- Edge density: 0.084.
- Average edge strength: 0.036.

Noise Residual Layer: Estimate edge/noise residual variance across a 4x4 tile grid.

- Noise tile variance mean: 986.6237.

- Noise tile inconsistency: 0.4813.

Compression And ELA Layer: Recompress image to JPEG quality 90 and measure error-level deltas plus JPEG markers.

- ELA normalized mean: 0.0.

- JPEG DQT marker count: 85.

Limitations

- Image authenticity cannot be proven from pixels alone.

- Social platforms often strip metadata, so missing EXIF is not proof of AI generation.

- No login scraping, private account access, face-search, or private identity inference was performed.

Technical Appendix

SHA-256: 7a3b93227153a932c461e92f7ec0a4ca573fced1cc960b5c8167793bd442ab96

Average hash: 3fdf7e2602d99006

Difference hash: 95692bab694cce63

Runtime: 85450.64 ms

This report is an evidence summary. It is not legal advice and does not certify authenticity.