

AI Deepfake Analyzer Report

The analysis is inconclusive.

Verdict: inconclusive with low confidence. Estimated AI probability is 34%; manipulation probability is 62%. Review the evidence layers before taking action.

Verdict	inconclusive
Confidence	low
AI Probability	34%
Manipulation Probability	62%
Detector Disagreement	53%

Decision Summary

The evidence does not cross a safe threshold; the report is preserving uncertainty instead of forcing a binary answer.

Primary drivers	Compression And ELA Layer: supports_real_or_camera_origin (AI 1%, manipulation 17%).
Counter-evidence	Source, Metadata, And Provenance: weak_anomaly (AI 50%, manipulation 62%).
Uncertainty	Confidence is low, so this should be treated as triage evidence rather than a final forensic certificate. Detector disagreement is elevated at 53%. No signed C2PA/content credential was available to independently confirm provenance. 1 layer(s) produced weak anomaly signals that need context.
What would help	Original, uncompressed media from the device or platform export. Public source URL, timestamps, captions, and account/page context. Signed C2PA/content credentials where available. Manual review when the result affects safety, reputation, or legal action.

Calibration Gate

Policy: victim_safe_calibrated_triage

Likely AI requires	Readable generative metadata or C2PA declaration,; or independent non-model fo
False-positive controls	Raw model scores are reliability-weighted before aggregation.; Detector disagree
Input quality risk	low

Mixture Of Experts

Expert	Opinion	Stance	Score
Visual Detector Ensemble	ensemble_inconclusive	neutral	30%
Portrait Specialist	portrait_gate_not_met	supports_ai	
Metadata And Provenance	no_generative_marker	neutral	50%
Forensic Residuals	neutral	neutral	17%
Input Quality Guard	low_quality_risk	neutral	18%

Safety Arbiter	inconclusive	neutral	34%
----------------	--------------	---------	-----

Evidence Layers

Input and Source - complete

- Image dimensions: 612 x 408.
- Input source: direct upload.

Metadata - complete

- EXIF metadata present.
- No GPS metadata was exposed.
- Software markers: XMP metadata mentions: photoshop.

Provenance - unavailable

- No c2patool/c2pa executable was found on PATH.

Hashes - complete

- SHA-256: 2ae007fbc16880208ba5fb699e52806cb0eadf41b12fc115cfd92b995b275fd8.
- Perceptual hash: 00307878fcfe70.

Compression and Noise - complete

- Manipulation heuristic score: 0.1726.
- Artificiality heuristic score: 0.0124.

Detector Ensemble - complete

- metadata_provenance: possibly_edited (ok).
- compression_noise_forensics: no_strong_forensic_signal (ok).
- hf:Ateeqq/ai-vs-human-image-detector: model_likely_human_or_real (ok).
- hf:dima806/ai_vs_real_image_detection: model_inconclusive (ok).
- hf:jacoballessio/ai-image-detect-distilled: model_inconclusive (ok).

Explainable AI Trace

- Final label is inconclusive with low confidence.
- Combined AI probability is 34%.
- Combined manipulation probability is 62%.
- Layer ledger: 0 support AI, 1 support real/camera-origin, 0 support manipulation, 1 weak anomalies, 8 neutral/inconclusive.

Region Evidence Map

This abstract map compares regions against the rest of the same image. It does not display the uploaded image and does not identify what caused a regional difference.

R1 C1 70% medium	R1 C2 23% low	R1 C3 58% medium	R1 C4 81% high
R2 C1 45% medium	R2 C2 100% high	R2 C3 22% low	R2 C4 81% high
R3 C1 51% medium	R3 C2 61% medium	R3 C3 30% low	R3 C4 52% medium
R4 C1 42% low	R4 C2 39% low	R4 C3 57% medium	R4 C4 67% medium

Analytical Layer Breakdown

Layer	Conclusion	AI	Manip.
Source, Metadata, And Provenance	weak_anomaly	50%	62%
Input Quality And Robustness	neutral_or_inconclusive	50%	18%
Visual Model Consensus Layer	neutral_or_inconclusive	30%	25%
Luminance Layer	neutral_or_inconclusive	36%	25%
Chroma And Color Layer	neutral_or_inconclusive	37%	25%
Edge And Geometry Layer	neutral_or_inconclusive	45%	45%
Noise Residual Layer	neutral_or_inconclusive	35%	45%
Compression And ELA Layer	supports_real_or_camera_origin	1%	17%
Frequency Spectrum Layer	neutral_or_inconclusive	58%	40%
Tile Region Anomaly Layer	neutral_or_inconclusive	43%	44%

Source, Metadata, And Provenance: EXIF/XMP/PNG text extraction plus optional C2PA manifest inspection.

- Generative markers: none.
- Editing markers: photoshop.

Input Quality And Robustness: Check resolution, file size, aspect ratio, entropy, and format conditions known to affect detector reliability.

- Robustness risk flags: moderate_resolution.
- Resolution: 612 x 408; shortest edge 408px.

Visual Model Consensus Layer: Run configured Hugging Face image-classification detectors and aggregate votes.

- 0/3 models voted AI-generated.
- 1/3 models voted real/human-origin.

Luminance Layer: Analyze grayscale entropy, contrast, clipping, and local brightness variation.

- Grayscale entropy: 7.716.
- Luminance contrast standard deviation: 0.213.

Chroma And Color Layer: Analyze RGB channel spread, channel correlation, and saturation distribution.

- Mean saturation: 0.391.
- Saturation variation: 0.277.

Edge And Geometry Layer: Run edge extraction and compare edge density/variance across tiles.

- Edge density: 0.249.
- Average edge strength: 0.117.

Noise Residual Layer: Estimate edge/noise residual variance across a 4x4 tile grid.

- Noise tile variance mean: 2995.8424.

- Noise tile inconsistency: 0.4543.

Compression And ELA Layer: Recompress image to JPEG quality 90 and measure error-level deltas plus JPEG markers.

- ELA normalized mean: 0.0.

- JPEG DQT marker count: 2.

Limitations

- Image authenticity cannot be proven from pixels alone.

- Social platforms often strip metadata, so missing EXIF is not proof of AI generation.

- No login scraping, private account access, face-search, or private identity inference was performed.

Technical Appendix

SHA-256: 2ae007fbc16880208ba5fb699e52806cb0eadf41b12fc115cfd92b995b275fd8

Average hash: 00307878fcfe70

Difference hash: 9e9f2f2c1f9b3759

Runtime: 2444.6 ms

This report is an evidence summary. It is not legal advice and does not certify authenticity.